

```

function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end

```

```

function GRAPH-SEARCH(problem, fringe) return a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE[INITIAL-STATE[problem]], fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    if STATE[node] is not in closed then
      add STATE[node] to closed
      for child-node in EXPAND(STATE[node], problem) do
        fringe ← INSERT(child-node, fringe)
      end
    end
  end

```

DFS, m level, b children of one node,

Time $O(b^m)$ Space $O(bm)$

Is it optimal? No, it finds the “leftmost” solution,

regardless of depth or cost

BFS time $O(b^b)$ s is the depth of solution

Space $O(b^b)$, ans it is complete, it is optimal if costs are all 1

UCS = uniform cost search

Is guaranteed to return an optimal path

Expand a cheapest node first, akes time $O(b^{C'/\epsilon})$, If that

solution costs C^* and arcs cost at least ϵ , then the

“effective depth” is roughly C^*/ϵ , space is the same,

Processes all nodes with cost less than cheapest solution!

Search Heuristics

Admissible: $0 \leq h(n) \leq h^*(n)$

$h^*(n)$ is the true cost to a nearest goal

Consistency: heuristic “arc” cost \leq actual cost for each arc $h(A) - h(C) \leq \text{cost}(A \text{ to } C)$

Greedy Search-best first search

Expand a mode that you think is closest to a goal state

A* Search

Admissible heuristics never underestimate? False 不会高估 cost

Uniform-cost =g(n), greedy-cost= h(n)

A* Search orders by the sum: f(n) = g(n) + h(n), stop:

when we de-queue a goal

The different between a* and ucs is that a* add the

heuristics which will intend to lead the agent to the goal with a purpose

Any complete search algorithm must have exponential (or worse) time complexity. True! (why?)

BFS always returns the optimal cost path. False (in general)

Before expanding a node, check to make sure its state has never been expanded before

If not new, skip it, if new add to closed set

Important: store the closed set as a set, not a list

Constraint Search problem

Backtracking Search: Depth-first search

with these two improvements is called

backtracking search (not the best name)

Idea 1: One variable at a time

Idea 2: Check constraints as you go

Improvement: Arc consistency & forward checking

```

function BACKTRACKING-SEARCH(csp) returns solution/failure
  return RECURSIVE-BACKTRACKING({}, csp)
function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment given CONSTRAINTS[csp] then
      add { var = value } to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result ≠ failure then return result
      remove { var = value } from assignment
  return failure

```

```

function AC-3(csp) returns the CSP, possibly with reduced domains
inputs: csp, a binary CSP with variables  $\{X_1, X_2, \dots, X_n\}$ 
local variables: queue, a queue of arcs, initially all the arcs in csp

while queue is not empty do
   $(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\text{queue})$ 
  if REMOVE-INCONSISTENT-VALUES( $X_i, X_j$ ) then
    for each  $X_k$  in NEIGHBORS[ $X_i$ ] do
      add  $(X_k, X_i)$  to queue

function REMOVE-INCONSISTENT-VALUES( $X_i, X_j$ ) returns true iff succeeds
  removed ← false
  for each  $x$  in DOMAIN[ $X_i$ ] do
    if no value  $y$  in DOMAIN[ $X_j$ ] allows  $(x, y)$  to satisfy the constraint  $X_i \leftrightarrow X_j$ 
      then delete  $x$  from DOMAIN[ $X_i$ ]; removed ← true
  return removed

```

Runtime: $O(n^2 d^3)$, can be reduced to $O(n^2 d^2)$

but detecting all possible future problems is NP-hard

ARC consistency- also called 2-consistency-a kind of binary consistency can affect any assigned variable in the

graph that has a path to the assigned variable

Limitation:1, can have one solution left. 2, have multiply

solution left. 3, none

Ordering: choose the MRV(Minimum(choose value)

remaining values), Choose the variable with the fewest

legal left values in its domain, also called most

constrained variable, fall-fast ordering

(which variable should be assigned next)

Ordering: Least Constraining Value

(choose node)

Given a choice of variable, choose the *least constraining value*

(in what order should its values be tired)

Iterative Improvement & local search: try incorrect solution and modify it to be correct.

Algorithm: while not solved:

randomly select any conflicted variable

Value selection: min-conflicts heuristic:

Choose a value that violates the fewest constraint

I.e., hill climb with $h(n)$ = total number of violated constraints

local search: improve a single option until you can’t make it better (no fringe!)

1,start whatever 2, repeat, move to the

best neighboring state 3, if no better

neighbors than current, return and quit

def max-value(state): def min-value(state):

initialize v = -∞

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$: $v = \min(v,$

$\text{value}(\text{successor}))$

return v **def value(state):**

if the state is a terminal state: return the state’s utility

if the next agent is MAX: return max-value(state)if the

next agent is MIN: return min-value(state) **time,**

space :DFS(without depth limited)

Alpha-Beta Implementation

α : MAX’s best option on path to root

β : MIN’s best option on path to root

def max-value(state, α, β):

initialize v = -∞

for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

if $v \geq \beta$ return v

$\alpha = \max(\alpha, v)$

return v

def min-value(state, α, β):

initialize v = +∞

for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$

if $v \leq \alpha$ return v

$\beta = \min(\beta, v)$

return v

Time complexity drops to $O(b^{m/2})$

Markov Decision Processes- **MDPs** are

non-deterministic search problems

One difference between MDP and

expectimax: rewards are smeared

throughout the tree(exit-max) instead of at the end.(MDP)

The value (utility) of a state s:

$V^*(s)$ = expected utility starting in s and

acting optimally

The value (utility) of a q-state (s,a):

$Q^*(s,a)$ = expected utility starting out

having taken action a from state s and

(thereafter) acting optimally (bellman

equation)

$V^*(s) = \max_a Q^*(s, a)$

$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$

$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$

T 是概率 V like MAX and Q like EXP. But

now incorporates rewards and discounts.

transition function

Optimal action: action with highest Q value

$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$

Complexity of each iteration: $O(S^2A)$

How do we know the V_k vectors are going

to converge?

Case 2: If the discount is less than 1

So as k increases, the values converge

$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$

Utilities for a Fixed Policy

calculate the V ’s for a fixed policy π (Policy

Evaluation)

$V_{k+1}^\pi(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k^\pi(s')]$

efficiency: $O(S^2)$ per iteration

problem:1, slow 2, the max at each state

rarely changes 3, the policy often

converges long before the values

policy iteration: Evaluation


$V_{k+1}^\pi(s) \leftarrow \sum_{s'} T(s, \pi_i(s), s') [R(s, \pi_i(s), s') + \gamma V_k^\pi(s')]$

improvement: For fixed values, get a

better policy using policy extraction

$\pi_{i+1}(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^\pi(s')]$

Example: Temporal Difference Learning (model base learning)



Value learning: watch some agent’s performance under a policy and see how well it does

Q-learning: watch yourself perform, possibly suboptimally, and still derive optimal policies/values/etc.

$V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + \alpha [R(s, \pi(s), s') + \gamma V^\pi(s')]$

if we want to turn values into a

(new) policy, we’re sunk:

$\pi(s) = \arg \max_a Q(s, a)$

$$Q(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V(s')]$$

so keep a running average

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + (\alpha) [r + \gamma \max_{a'} Q(s', a')]$$

approximate Q-learning

$$\text{difference} = [r + \gamma \max_{a'} Q(s', a')] - Q(s, a)$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha [\text{difference}]$$


$$w_i \leftarrow w_i + \alpha [\text{difference}] f_i(s, a)$$

example:

$$Q(s, a) = 4.0 f_{DOT}(s, a) - 1.0 f_{GST}(s, a)$$

wdot = 4, wgst = -1

Example: Q-Pacman



conditional P: $P(a|b) = P(a, b) / P(b)$

product rule: $P(y)P(x|y) = P(x, y)$

$P(W|T = c)$ w sun; p = 0.4

w rain: p = 0.6

Inference by Enumeration: Worst-case time

complexity $O(d^n)$, Space complexity $O(d^n)$ to store the joint distribution

Chain rule:

$$P(x_1, x_2, x_3) = P(x_1)P(x_2|x_1)P(x_3|x_1, x_2)$$

Baye's rule: $P(x|y) = \frac{P(y|x)P(x)}{P(y)} = P(x, y) / P(y)$

$$P(y) = \sum_x P(x, y) = P(+m)P(y|m) + P(-m)P(y|-m)$$

$$= \sum_{x_{t-1}} P(x_t | x_{t-1}) P(x_{t-1})$$

Sum = 1

Using baye's rule: $P(G) = 0.11$ for all is uniform

$P(r=yellow | ghost \text{ at } (1,1)) = 0.1$

So $P(g|r) = P(g)P(r|g) = 0.1 * 0.11 = 0.011$

X is conditionally independent of Y given Z:

$$\forall x, y, z : P(x, y|z) = P(x|z)P(y|z)$$

$$\forall x, y, z : P(x|z, y) = P(x|z)$$

Joint Distribution of a Markov Model

核心: now only related to past, future only related to now

$$P(X_1, X_2, X_3, X_4) = P(X_1)P(X_2|X_1)P(X_3|X_2)P(X_4|X_3)$$

$$X_3 \perp\!\!\!\perp X_1 \mid X_2 \quad X_4 \perp\!\!\!\perp X_1, X_2 \mid X_3$$

Stationary Distributions: 最终 convergent :

$$P_\infty(X) = P_{\infty+1}(X) = \sum_x P(X|x)P_\infty(x)$$

HMM: evidence 're only conditionally independent given the hidden state

Quiz: does this mean that evidence variables are guaranteed to be independent?

No, they tend to correlated by the hidden state!

Joint distribution of a HMM:

$$P(X_1, E_1, X_2, E_2, X_3, E_3) = P(X_1)P(E_1|X_1)P(X_2|X_1)P(E_2|X_2)P(X_3|X_2)P(E_3|X_3)$$

$$P(X_1, E_1, \dots, X_T, E_T) = P(X_1)P(E_1|X_1) \prod_{t=2}^T P(X_t|X_{t-1})P(E_t|X_t)$$

inference: base case:

$$P(X_1|e_1)$$

$$\begin{aligned} P(x_1|e_1) &= P(x_1, e_1) / P(e_1) \\ &\propto_{X_1} P(x_1, e_1) \\ &= P(x_1)P(e_1|x_1) \end{aligned}$$

if we have e_1 to e_t and the $B(X_t) = P(X_t|e_{1:t})$, we reason

$$= \sum_{x_t} P(X_{t+1}, x_t|e_{1:t})$$

about : $P(X_{t+1}|e_{1:t}) =$

$$= \sum_{x_t} P(X_{t+1}|x_t)P(x_t|e_{1:t})$$

or we can say:

$$B'(X_{t+1}) = \sum_{x_t} P(X'|x_t)B(x_t)$$

this is the guess part

the following is the observation part:

we get the believe for X_{t+1} :

$$B'(X_{t+1}) = P(X_{t+1}|e_{1:t})$$

then the evidence comes in

$$P(X_{t+1}|e_{1:t+1}) =$$

$$P(X_{t+1}, e_{t+1}|e_{1:t}) / P(e_{t+1}|e_{1:t})$$

$$\propto_{X_{t+1}} P(X_{t+1}, e_{t+1}|e_{1:t})$$

$$= P(e_{t+1}|e_{1:t}, X_{t+1})P(X_{t+1}|e_{1:t})$$

$$= P(e_{t+1}|X_{t+1})P(X_{t+1}|e_{1:t})$$

Unlike passage of time, we have to renormalize

Basic idea: beliefs "reweighted" by likelihood of evidence, and compactly:

$$B(X_{t+1}) \propto_{X_{t+1}} P(e_{t+1}|X_{t+1})B'(X_{t+1})$$

the forward algorithm

$$\text{we want to know } B_t(X) = P(X_t|e_{1:t})$$

we can derive the following updates:

$$P(x_t|e_{1:t}) \propto_X P(x_t, e_{1:t})$$

$$= \sum_{x_{t-1}} P(x_{t-1}, x_t, e_{1:t})$$

$$= \sum_{x_{t-1}} P(x_{t-1}, e_{1:t-1})P(x_t|x_{t-1})P(e_t|x_t)$$

$$= P(e_t|x_t) \sum_{x_{t-1}} P(x_t|x_{t-1})P(x_{t-1}, e_{1:t-1})$$

SUMMARY:

$$P(x_t|e_{1:t-1}) = \sum_{x_{t-1}} P(x_{t-1}, e_{1:t-1}) \cdot P(x_t|x_{t-1})$$

We reason that

Get evidence: we get: (update)

$$P(x_t|e_{1:t}) \propto_X P(x_t|e_{1:t-1}) \cdot P(e_t|x_t)$$

Particle filter:

Step 1: base on the evidence-> guess what will happen

$$x' = \text{sample}(P(X'|x))$$

Step 2: observe: $w(x) = P(e|x)$

$B(X) \propto P(e|X)B'(X)$ As before, the probabilities don't sum to one, since all have been down weighted

Step 3: resample: normalize

HMMS: most likely explanation: forward algorithm (Sum)

$$\arg \max_{x_{1:t}} P(x_{1:t}|e_{1:t}) \quad f_t[x_t] = P(x_t, e_{1:t})$$

$$Viterbi \text{ Algorithm}(\max): m_t[x_t] = \max_{x_{1:t-1}} P(x_{1:t-1}, x_t, e_{1:t})$$

$$= P(e_t|x_t) \max_{x_{t-1}} P(x_t|x_{t-1})m_{t-1}[x_{t-1}]$$

which state sequence $X_{1:T}$ is most likely given the

evidence $e_{1:T}$:

$$x_{1:T}^* = \arg \max_{x_{1:T}} P(x_{1:T}|e_{1:T}) = \arg \max_{x_{1:T}} P(x_{1:T}, e_{1:T})$$

Bayes Net: to see what probability a BN gives to a full assignment: 假设 x 只与 parent (x) 有关

$$P(x_1, x_2, \dots, x_n) = \prod_{i=1}^n P(x_i | \text{parents}(X_i))$$

same :

$$P(x_i|x_1, \dots, x_{i-1}) = P(x_i | \text{parents}(X_i))$$

How big is an N-node net if nodes have up to k parents?

$$O(N * 2^{k+1})$$

General question: in a given BN, are two variables independent (given evidence)?

X->Y->Z, Guaranteed X independent of Z? No!

Guaranteed X independent of Z given Y? yes!

Y->X, Y->Z. Guaranteed X independent of Z? No!

Guaranteed X independent of Z given Y? yes!

Evidence along the chain "blocks" the influence

X->Z, Y->Z. Are X and Y independent? Yes

Are X and Y independent given Z? NO: seeing traffic puts the rain and the ballgame in competition as explanation.

Active means no independent (dependent)

A trail $X_1 \dots X_n$ is active if: it has no v-structures $X_{i-1} \rightarrow X_i \leftarrow X_{i+1}$

A trial $X_1 \dots X_n$ is active given Z if :

1, for any v-structure $X_{i-1} \rightarrow X_i \leftarrow X_{i+1}$ we have that X_i or one of its descendants $\in Z$, no other X_i is in Z


如果两个节点间没有 Active trial, 则这两个节点是 d-separated 的 也就是独立的! Independence

$X_i \perp\!\!\!\perp X_j | \{X_{k_1}, \dots, X_{k_n}\}$? Given means Block (DARKED)

Check all (undirected!) paths between X_i and X_j

If one or more active path, then independence not guaranteed

If all paths are inactive, then independence is guaranteed



Decision Networks: choose the action which maximizes the expected utility given the evidence

= EU(ACTION) =

$$EU(\text{leave}) = \sum_w P(w)U(\text{leave}, w)$$

MEU(X)=max(EU for all) 在你获得新的 evidence 之后, you update you believe (probability), but the value of Utility doesn't change! Evidence kinds of forecast:

P(rain | forecast = bad) = 0.66, P(sun | forecast = bad) =

0.34, recalculate the EU, and MEU

Value of information kinds of evidence =

MEU (after we get the information)- MEU(before we get the information)

$$VPI(E'|e) = \left(\sum_{e'} P(e'|e)MEU(e, e') \right) - MEU(e)$$

$$\forall E', e : VPI(E'|e) \geq 0$$

1, nonadditive:

$$VPI(E_j, E_k|e) \neq VPI(E_j|e) + VPI(E_k|e)$$

$$VPI(E_j, E_k|e) = VPI(E_j|e) + VPI(E_k|e, E_j)$$

$$= VPI(E_k|e) + VPI(E_j|e, E_k)$$

If Parents(U) $\cap Z \mid$ CurrentEvidence

Then $VPI(Z \mid \text{CurrentEvidence}) = 0$